

Time series

Preliminary definitions

Time series definition [2]

Informal definition

A time-series is a set of observation x_t each one being recorder at a specific time t .

Formal definition

A time series model for the observed data x_t is a specification of the joint distribution (or possible only the means covariance) of a sequence of random variable X_t of which x_t is postulated to be a realization

A binary process

Consider the sequence of iid random variables, with $P[X_t = 1] = p$ and $P[X_t = -1] = 1 - p$

Random walk

The random walk is obtained by cumulatively summing iid random variables. Thus a random walk with zero mean is obtained by defining

Stationarity, autocovariance and autocorrelation[2]

Mean Function

Let X_t be a time series with $E(x_t^2) < \infty$. The mean function of X_t is $\mu_X(t) = E(X_t)$. The covariance function of X_t is $\gamma_X(r, s) = \text{Cov}(X_r, X_s) = E[(X_r - \mu_X(r))(X_s - \mu_X(s))]$ $\forall r, s$

Weakly stationary TS

X_t is weakly stationary if i) $\mu_X(t)$ is independent from time t and ii) $\gamma_X(t+h)$ is independent of $t \forall h$

Autocovariance function

At lag h the auto-covariance function is defined as $\gamma_X(h) = \text{Cov}(X_{t+h}, X_t)$

Autocorrelation function

At lag h the autocorrelation function is defined as

$$\rho(h)_X = \frac{\gamma_X(h)}{\gamma_X(0)} = \text{Cor}(X_{t+h}, X_t)$$

Stationarity, autocovariance and autocorrelation [2]

Figure 1-12
200 simulated values of IID
 $N(0,1)$ noise







Figure 1-13
The sample autocorrelation
function for the data of
Figure 1-12 showing the
bounds $\pm 1.96/\sqrt{n}$



Stationarity, autocovariance and autocorrelation [4]



(a)



(b)

Definition [2]

Linear process

Time series X_t is a linear process if it has the representation

$$X_t = \sum_{j=-\infty}^{\infty} \psi_j Z_{t-j}$$

for all t , where $Z_t \approx \text{WN}(0, \sigma^2)$ and ψ_j is a sequence of constants with $\sum_{j=-\infty}^{\infty} \psi_j < \infty$. In terms of backward shift operator B $BX_t = X_{t-1}$ we have $X_t = \phi(B)Z_t$. Therefore the previous definition can be rewritten as $X_t = \phi(B)Z_t$ in which $\phi(B)$ can be thought as a linear filter that when applied to the white noise input series Z_t produces the output X_t .

Definition [2]

Linear process

The time series X_t is an ARMA(1,1) process if its stationary and satisfiy (for every t)

$$X_t - \phi X_{t-1} = Z_t + \theta Z_{t-1}$$

where $Z_t \approx \text{WN}(0, \phi^2)$ and $\phi + \theta \neq 0$ or in terms of filters ϕ and θ

$$\phi(B)X_t = \theta(B)Z_t$$

Remarks

- A stationary solution of the ARMA(1,1) equation exists if and only if $\phi \neq \pm 1$
- If $|\phi| < 1$, then the unique stationary solution is given by $X_t = Z_t + (\phi + \theta) \sum_{j=1}^{\infty} \phi^{j-1} Z_{t-j}$. In this case we say that X_t is causal or a causal function of Z_t or a causal function of Z_t since X_t can be expressed in terms of the current and past values $Z_s, s \leq t$
- If $|\phi| > 1$, then the unique stationary solution is given by $X_t = -\theta\phi^{-1}Z_t - (\phi + \theta) \sum_{j=1}^{\infty} \phi^{-j-1}Z_{t-j}$. The solution is non-causal, since X_t is then a functional of $Z_s, s \geq t$

Wold decomposition [2]

Prediction operator based on the infinite past $X_t, -\infty < t < n$

$$\tilde{P}_n X_{n+h} = \lim_{m \rightarrow -\infty} P_{m,n} X_{n+h}$$

Wold decomposition $X_t, -\infty < t < n$

X_t is a non-deterministic stationary time series, then

$$X_t = \sum_{j=0}^{\infty} \psi_j Z_{t-j} + V_t$$

where V_t is deterministic

Definition [2]

ARMA (p,q)

X_t is an ARMA(p,q) process if X_t is stationary and if for every t,

$$X_t - \phi_1 X_{t-1} - \dots - \phi_p X_{t-p} = Z_t + \theta_1 Z_{t-1} + \dots + \theta_q Z_{t-q}$$

where $Z_t \approx WN(0, \sigma^2)$ and the polynomials $(1 - \phi_1 z - \dots - \phi_p z^p)$ and $(1 + \theta_1 z + \dots + \theta_q z^q)$ have no common factor.

Definition [2]

Causality

An ARMA(p,q) process X_t is causal, or a causal function of Z_t if there exist constants ϕ_j such that $\sum_{j=0}^{\infty} |\phi_j| < \infty$ and

$$X_t = \sum_{j=0}^{\infty} \sum \phi_j Z_{t-j} \quad \forall \quad t$$

this is equivalent to the condition

$$\phi(z) = 1 - \phi_1 z - \dots - \phi_p z^p \neq 0 \quad |z| < 1$$

Spectral Analysis

Spectral Analysis [2]

Spectral density

Given a zero mean stationary time series X_t with autocovariance function $\gamma(\cdot)$ satisfying $\sum_{h=-\infty}^{\infty} |\gamma(h)| < \infty$, the spectral density of X_t is the function defined by

$$f(\lambda) = \frac{1}{2\pi} \sum e^{-ih\lambda} \gamma(h) \quad -\infty < \lambda < \infty$$

with the condition that

$$f(\lambda) \geq 0 \quad \forall \lambda$$

$$\gamma(h) = \int_{-\pi}^{\pi} e^{ih\lambda} f(\lambda) d\lambda \quad \forall h$$

Spectral Analysis [2]

Spectral Representation of The ACVF

A function $\gamma(\cdot)$ defined on the integers is the ACVF of a stationary time series if and only if there exists a right-continuous, non-decreasing, bounded function F on $[-\pi, \pi]$ with $F(-\pi) = 0$ such that:

$$\gamma(h) = \int_{(-\pi, \pi]} e^{ih\lambda} dF(\lambda)$$

for all integers h . F is a generalized distribution function that is called the spectral distribution function of $\gamma(\cdot)$

Spectral Analysis [2]

Figure 4-1
A sample path of size
100 from the time series
in Example 4.1.2





Figure 4-2
The spectral distribution
function $F(\lambda)$, $-\pi \leq \lambda \leq \pi$,
of the time series
in Example 4.1.2



Time-Invariant Linear Filters [2]

Linear process

The process Y_t is the output of a linear filter $C = \{c_{t,k}, t, k = 0 \pm 1, \dots\}$ applied to an input process X_t if $Y_t = \sum_{k=-\infty}^{\infty} c_{t,k} X_k \quad t = 0, \pm 1, \dots$

Time invariant

The filter is said to be time-invariant if the weights $c_{t,t-k}$ are independent of t e.g. $c_{t,t-k} = \phi_k$

$$Y_t = \sum_{k=-\infty}^{\infty} \phi_{t,k} X_{t-k}$$

$$Y_{t-s} = \sum_{k=-\infty}^{\infty} \phi_{t,k} X_{t-s-k}$$

The TLF ϕ is to be causal if $\phi_j = 0$ for $j < 0$

Time-Invariant Linear Filters [2]

Transfer function

Let X_t be a stationary time series with mean zero and spectral density $f_X(\lambda)$. Suppose that $\Phi = \{\phi_j, j = 0, \pm 1, \dots\}$ is an absolutely summable TLF. Then the time series

$$Y_t = \sum_{j=-\infty}^{\infty} \phi_j X_{t-j}$$

is stationary with mean zero and spectral density

$$f_Y(\lambda) = |\Psi(e^{-i\lambda})|^2 f_X(\lambda) = \Psi(e^{-i\lambda}) \Psi(e^{i\lambda}) f_X(\lambda)$$

where $\Psi(e^{-i\lambda}) = \sum_{j=-\infty}^{\infty} \phi_j e^{-ij\lambda}$ where $\Psi(e^{-i})$ is called the transfer function of the filter, and the squared modulus $|\Psi(e^{-i})|$ is referred to as the power transfer function of the filter.

Time-Invariant Linear Filters [2]

Simple moving average

$$Y_t = \frac{1}{2q+1} \sum_{|j| \leq q} X_{t-j}$$

where $\psi = (2q+1)^{-1}$, $j = -q, \dots, q$ and ψ_j




Figure 4-12
The transfer function $D_{10}(\lambda)$ for the simple moving-average filter

Time-Invariant Linear Filters [2]

Gibbs phenomenon

The poor approximation in the neighbourhood of cut-off frequency (ω_c)




Figure 4-13

The transfer function for the ideal low-pass filter and truncated Fourier approximations $\Psi^{(q)}$ for $q = 2, 10$

Estimation

Estimation [2]

Gaussian Likelihood

$$L(\hat{\boldsymbol{\phi}}, \hat{\boldsymbol{\theta}}, \sigma^2) = \frac{1}{\sqrt{(2\pi\sigma)^n r_0 \dots r_{n-1}}} \exp\left(-\frac{1}{2\sigma^2} \sum \frac{(X_j - \hat{X}_j)^2}{r_{j-1}}\right)$$

Maximum Likelihood Estimators

Differentiating $\ln L(\boldsymbol{\phi}, \boldsymbol{\theta}, \sigma^2)$ partially with respect to σ^2 and noting that \hat{X}_j and r_j are independent of σ

$$\hat{\sigma}^2 = n^{-1} S(\hat{\boldsymbol{\phi}}, \hat{\boldsymbol{\theta}})$$

$$S(\hat{\boldsymbol{\phi}}, \hat{\boldsymbol{\theta}}) = \sum_{j=1}^n (X_j - \hat{X}_j)^2 / r_{j-1}$$

where $\hat{\boldsymbol{\phi}}, \hat{\boldsymbol{\theta}}$ are the values of $\boldsymbol{\phi}, \boldsymbol{\theta}$ that minimize:

$$\mathcal{L}(\boldsymbol{\phi}, \boldsymbol{\theta}) = \ln(n^{-1} S(\boldsymbol{\phi}, \boldsymbol{\theta})) + n^{-1} \sum_{j=1}^n \ln r_{j-1}$$

Estimation [2]

Akaike information criterion

$$AICC = -2 \ln L(\hat{\Phi}, \hat{\theta}, S((\phi_p, \theta_q)/n) + 2(p + q + 1)n/(n - p - q - 2)$$

Residual

$$\hat{W}_t = (X_t - \hat{X}_t \phi, \theta) / (r_{t-1}(\hat{\Phi}, \hat{\theta}))^{1/2} \quad t = 1, \dots, n$$

This should have properties similar to the white noise sequence, thus one can define the rescaled residuals

$$\hat{R}_t = \hat{W}_t / \hat{\sigma}$$

$$\hat{\sigma} = \sqrt{\left(\sum_{t=1}^n \hat{W}_t^2 \right) / n}$$

Residuals [2]








Figure 5-5
The rescaled residuals after
fitting the ARMA(1,1) model
of Example 5.2.5 to the lake
data

Nonstationary time series

Residuals [2, 1]



Akaike information criterion

If d is a non-negative integer, then X_t is an ARIMA(p, d, q) process if

$$Y_t = (1 - B)^d X_t$$

is causal ARMA(p, q) process $BX_t = X_{t-1}$. Note that $\nabla X_t = X_t - X_{t-1}$

Multivariate time series

Multivariate time series [2]

Akaike information criterion

$\{\mathbf{X}_t\}$ is an ARMA (p,q) process if $\{\mathbf{X}_t\}$ is stationary and if for every t,

$$\mathbf{X}_t - \phi_1 \mathbf{X}_{t-1} - \dots - \phi_p \mathbf{X}_{t-p} = Z_t + \Theta_1 Z_{t-1} + \dots + \Theta_q Z_{t-q}$$

where $Z_t \approx \text{WN}(0, \Sigma)$


Vector Autoregression [3]

VAR(1)

$$\begin{pmatrix} y_{1,t} \\ y_{2,t} \end{pmatrix} = \begin{pmatrix} c_1 \\ c_2 \end{pmatrix} + \begin{pmatrix} a_{1,1} & a_{1,2} \\ a_{2,1} & a_{2,2} \end{pmatrix} \begin{pmatrix} y_{1,t-1} \\ y_{2,t-1} \end{pmatrix} + \begin{pmatrix} e_{1,t} \\ e_{2,t} \end{pmatrix}$$

Granger causality test (hamilton2020time)


y_2 cause (in Granger sense) y_1 if the coefficient $a_{1,2}$ is significantly not equal to 0



Vector Autoregression [3]

Impulse response function


y_2 cause (in Granger sense) y_1 if the coefficient $a_{1,2}$ is significantly not equal to 0



Space State Models [5]

Definition

A state-space model (SSM) is a partially observed Markov model in which the hidden state z_t evolves over time according to a Markov process and each hidden state generates some observations y_t at each step. The main goal is to infer the hidden states given the observations (and also to predict the future observations)



Space State Models [5]

Non linear dynamical system

A state-space model (SSM) can be represented as a stochastic discrete time nonlinear dynamical system of the form

$$\begin{aligned}z &= f(\mathbf{z}_{t-1}, \mathbf{u}_t, \mathbf{q}_t) \\ y_t &= h(\mathbf{z}_t, \mathbf{u}_t, \mathbf{y}_{1:t-1}, \mathbf{r}_t)\end{aligned}$$

where $\mathbf{z}_t \in \mathbb{R}^N$ are the hidden states, $\mathbf{u}_t \in \mathbb{R}^N$ are the optional observed inputs, $\mathbf{y}_t \in \mathbb{R}^N$ are observed output and \mathbf{f} is the transition function, \mathbf{q}_t is the process noise, \mathbf{h} is the observation function and \mathbf{r}_t is the observation noise. The transition model and the observational model are


$$\begin{aligned}p(\mathbf{z}_t | \mathbf{z}_{t-1}, \mathbf{u}_t) &= p(\mathbf{z}_t | \mathbf{f}(\mathbf{z}_{t-1}, \mathbf{u}_t)) \\ p(\mathbf{y}_t | \mathbf{z}_t, \mathbf{u}_t, \mathbf{y}_{1:t-1}) &= p(\mathbf{y}_t | \mathbf{h}(\mathbf{z}_t, \mathbf{u}_t, \mathbf{y}_{1:t-1}))\end{aligned}$$

Space State Models [5]


Non linear dynamical system

Hidden Markov Model (HMM) \rightarrow a SSM in which the hidden states are discrete, thus $z_t \in 1, \dots, K$

Time series segmentation



(a)



(b)

Figure 29.6: (a) A sample time series dataset of counts. (b) A segmentation of this data using a 4 state HMM. Generated by `poisson_hmm_changepoint.py`.




Figure 29.7: Segmentation of the time series using HMMs with 1-6 states. Generated by `poisson_hmm_changepoint.py`.

Space State Models [5]

Time series segmentation

We want to segment a time series into different regimes, each of which correspond to a different statistical distribution. In particular we would like to segment this data stream in to K different regimes or states, each of which is associated with a Poisson observation model with rate λ_k :

$$p(y_t | z_t = k) = \text{Poi}(y_t | \lambda_k)$$

where a uniform prior over the initial states was considered. The transition matrix will be:

$$z_1 \approx \text{Categorical} \left(\left\{ \frac{1}{4}, \frac{1}{4}, \frac{1}{4}, \frac{1}{4} \right\} \right)$$

$$z_t | z_{t-1} \approx \text{Categorical} \left(\left\{ \begin{pmatrix} p & z_t = z_{t-1} \\ \frac{1-p}{4-1} & \text{else} \end{pmatrix} \right\} \right)$$

$$p(y_{1:T} | K) \approx \max_{\lambda} \sum p(y_{1:T}, z_{1:T} | \lambda, K)$$

Structural time series models [5]

Structural time series (STS)

Defined in terms of linear-Gaussian SSMs. Differently from ARMA method, they have much more flexibility: one can create non-linear, non-Gaussian and even hierarchical extension. Represent the observed scalar time series as a sum of C individual components

$$f(t) = f_1(t) + f_2(t) + \dots + f_C(t) + \epsilon_t$$

Each single component (latent process) $f_c(t)$ is modeled by a linear Gaussian state-space model which is also called dynamic linear model. Since these are linear, one can combine in to a single state-space model.

$$p(z_t|z_{t-1}, \theta) = \mathcal{N}(z_t|Fz_{t-1}, Q)$$

$$p(y_t|z_t, \theta) = \mathcal{N}(y_t|Hz_t + \beta^T u_t, \sigma_y^2)$$

where F and Q are block structure matrices, with one block per component. The vector H then adds up all the relevant pieces from each component to generate the overall mean.

Structural time series models [5]

Local level model

The observation $y_t \in \mathbb{R}^N$ are generated by a Gaussian with (latent) mean μ_t , which evolves over over time according to a random walk

$$\begin{aligned}y_t &= \mu_t + \epsilon_{y,t} & \epsilon_{y,t} &\approx \mathcal{N}(0, \sigma_y^2) \\ \mu_t &= \mu_{t-1} + \epsilon_{t-1\mu,t} & \epsilon_{\mu,t} &\approx \mathcal{N}(0, \sigma_\mu^2)\end{aligned}$$

One can also assume that $\mu_1 \in \mathcal{N}(0, \sigma_\mu^2)$. Thus the latent mean at any future step has distribution $\mu_t \in \mathcal{N}(0, t\sigma_\mu^2)$ so the variance grows with time. One can also use an autoregressive model in which $\mu_t = \rho\mu_{t-1} + \epsilon_{\mu,t}$ where $|\rho| < 1$. In this case we have $\mu_\infty \approx \mathcal{N}(0, \frac{\sigma_\mu^2}{1-\rho^2})$: thus the uncertainty grows to finite asymptote instead of undoubtedly

Bibliography I

- [1] <https://www.semanticscholar.org/paper/NON-STATIONARY-TIME-SERIES-AND-UNIT-ROOT-TESTING-Nielsen/2d0e62db75bdeafd2277fab1039f4866aef642b7>.
- [2] Peter J Brockwell e Richard A Davis. *Introduction to time series and forecasting*. Springer, 2002.
- [3] James D Hamilton. *Time series analysis*. Princeton university press, 2020.
- [4] Jianhua Hao e Fangai Liu. “Improving long-term multivariate time series forecasting with a seasonal-trend decomposition-based 2-dimensional temporal convolution dense network”. In: *Scientific Reports* 14.1 (2024), p. 1689.
- [5] Kevin P. Murphy. *Probabilistic Machine Learning: Advanced Topics*. MIT Press, 2023. URL: <http://probml.github.io/book2>.